


+

# Machine Learning and Data Mining



## Support Vector Machines

Prof. Alexander Ihler



# Linear Classifiers

- Which decision boundary is “better”?
  - Both have zero training error (perfect training accuracy)
  - But, one of them seems intuitively better...
- How can we quantify “better”,  
and learn the “best” parameter settings?



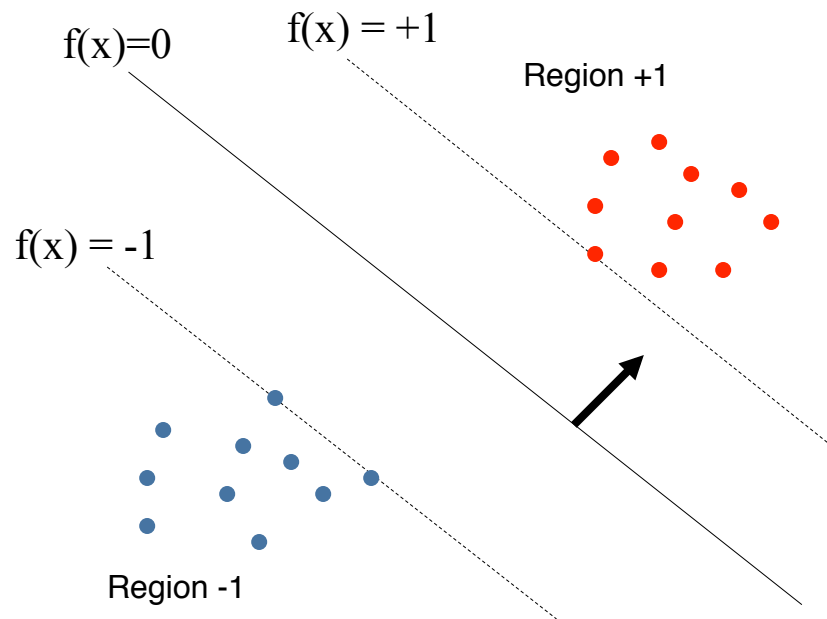
# One possible answer...

- Maybe we want to maximize our “margin”
- To optimize, relate to model parameters
- Remove “scale invariance”
  - Define class +1 in some region, class –1 in another
  - Make those regions as far apart as possible

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$



$$b + w_1 x_1 + w_2 x_2 + \dots$$



We could define such a function:

$$f(x) = w^*x' + b$$

$$f(x) > +1 \text{ in region +1}$$


$$f(x) < -1 \text{ in region -1}$$

Passes through zero in center...

**“Support vectors” – data points on margin**


# Computing the margin width

- Vector  $\underline{w}=[w_1 \ w_2 \ \dots]$  is perpendicular to the boundaries (why?)
- $w x + b = 0$  &  $w x' + b = 0 \Rightarrow w (x'-x) = 0$  : orthogonal



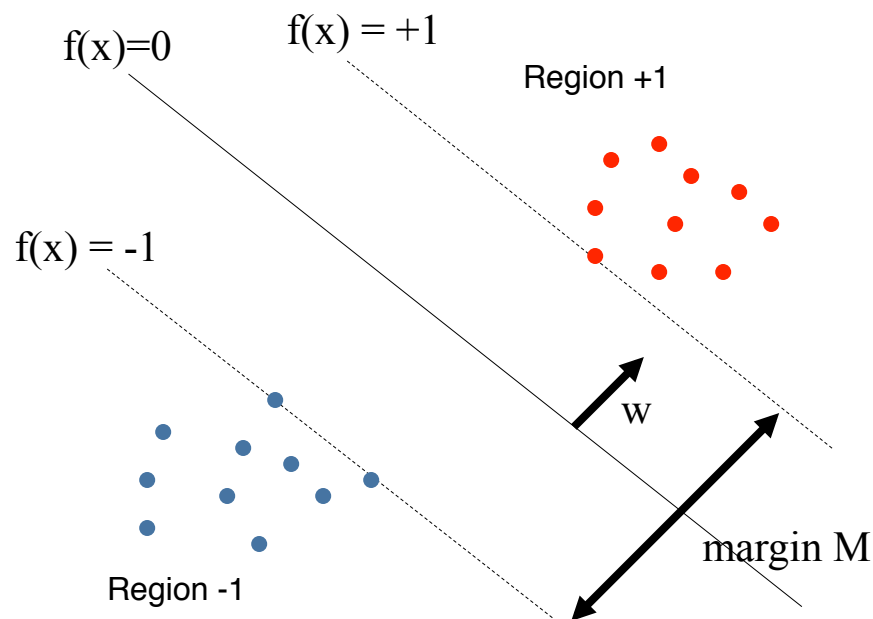
# Computing the margin width

- Vector  $\underline{w} = [w_1 \ w_2 \ \dots]$  is perpendicular to the boundaries
- Choose  $\underline{x}^-$  st  $f(\underline{x}^-) = -1$ ; let  $\underline{x}^+$  be the closest point with  $f(\underline{x}^+) = +1$ 
  - $\underline{x}^+ = \underline{x}^- + r * \underline{w}$  (why?)
- Closest two points on the margin also satisfy
$$w \cdot x^- + b = -1 \qquad w \cdot x^+ + b = +1$$



# Computing the margin width

- Vector  $\underline{w} = [w_1 \ w_2 \ \dots]$  is perpendicular to the boundaries
- Choose  $\underline{x}^-$  st  $f(\underline{x}^-) = -1$ ; let  $\underline{x}^+$  be the closest point with  $f(\underline{x}^+) = +1$ 
  - $\underline{x}^+ = \underline{x}^- + r * \underline{w}$
- Closest two points on the margin also satisfy
$$w \cdot x^- + b = -1 \qquad w \cdot x^+ + b = +1$$



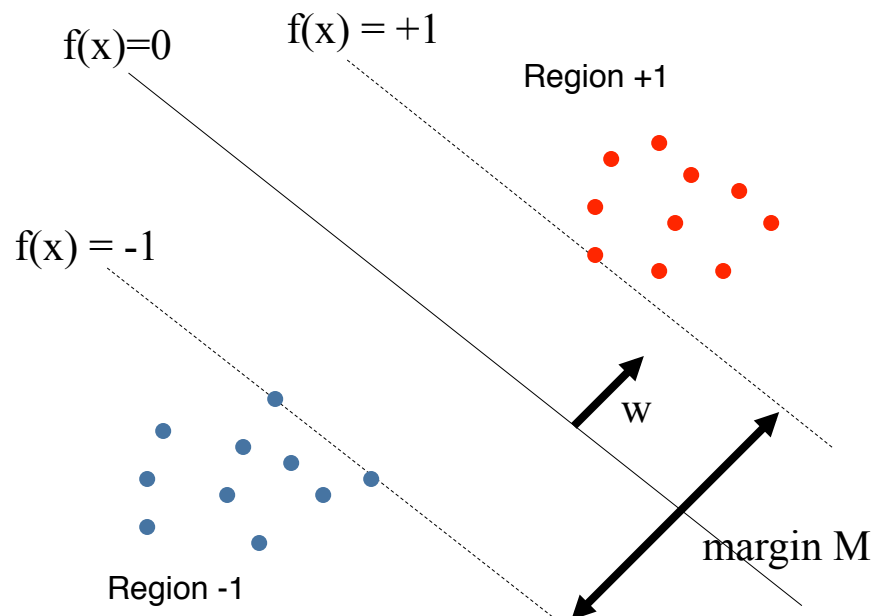
$$\begin{aligned} w \cdot (x^- + rw) + b &= +1 \\ \Rightarrow r\|w\|^2 + w \cdot x^- + b &= +1 \\ \Rightarrow r\|w\|^2 - 1 &= +1 \\ \Rightarrow r &= \frac{2}{\|w\|^2} \end{aligned}$$

$$\begin{aligned} M &= \|x^+ - x^-\| = \|rw\| \\ &= \frac{2}{\|w\|^2} \|w\| = \frac{2}{\sqrt{w^T w}} \end{aligned}$$

# Maximum margin classifier

- Constrained optimization
  - Get all data points correct
  - Maximize the margin

This is an example of a quadratic program:  
quadratic cost function, linear constraints



$$w^* = \arg \max_w \frac{2}{\sqrt{w^T w}}$$

such that “all data on the correct side of the margin”

## Primal problem:

$$w^* = \arg \min_w \sum_j w_j^2$$

s.t.

$$y^{(i)} = +1 \Rightarrow w \cdot x^{(i)} + b \geq +1$$

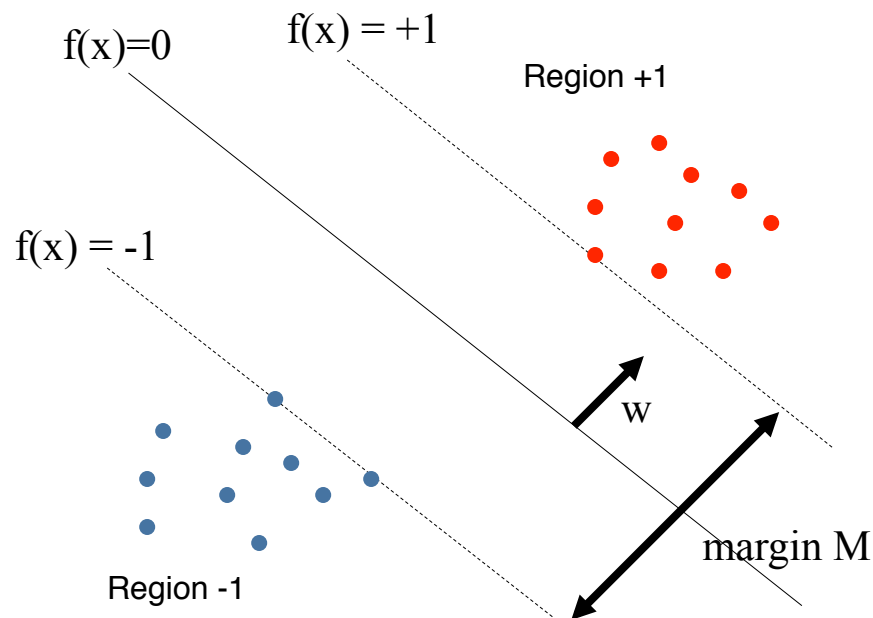
$$y^{(i)} = -1 \Rightarrow w \cdot x^{(i)} + b \leq -1$$

( $m$  constraints)

# Maximum margin classifier

- Constrained optimization
  - Get all data points correct
  - Maximize the margin

This is an example of a quadratic program:  
quadratic cost function, linear constraints



$$w^* = \arg \max_w \frac{2}{\sqrt{w^T w}}$$

such that “all data on the correct side of the margin”

**Primal problem:**

$$w^* = \arg \min_w \sum_j w_j^2$$

s.t.



$$y^{(i)} (w \cdot x^{(i)} + b) \geq +1$$

( $m$  constraints)





# A 1D Example

- Suppose we have three data points  
 $x = -3, y = -1$   
 $x = -1, y = -1$   
 $x = 2, y = 1$
- Many separating perceptrons,  $T[ax+b]$ 
  - Anything with  $ax+b = 0$  between -1 and 2
- We can write the margin constraints
$$\begin{aligned} a(-3) + b &< -1 && \Rightarrow b < 3a - 1 \\ a(-1) + b &< -1 && \Rightarrow b < a - 1 \\ a(2) + b &> +1 && \Rightarrow b > -2a + 1 \end{aligned}$$





# A 1D Example

- Suppose we have three data points  
 $x = -3, y = -1$   
 $x = -1, y = -1$   
 $x = 1, y = 1$
- Many separating perceptrons,  $T[ax+b]$ 
  - Anything with  $ax+b = 0$  between -1 and 2
- We can write the margin constraints
$$a(-3) + b < -1 \quad \Rightarrow b < 3a - 1$$
$$a(-1) + b < -1 \quad \Rightarrow b < a - 1$$
$$a(2) + b > +1 \quad \Rightarrow b > -2a + 1$$
- Ex:  $a = 1, b = 0$



# A 1D Example

- Suppose we have three data points  
 $x = -3, y = -1$   
 $x = -1, y = -1$   
 $x = 1, y = 1$
- Many separating perceptrons,  $T[ax+b]$ 
  - Anything with  $ax+b = 0$  between -1 and 2
- We can write the margin constraints  
 $a(-3) + b < -1 \Rightarrow b < 3a - 1$   
 $a(-1) + b < -1 \Rightarrow b < a - 1$   
 $a(2) + b > +1 \Rightarrow b > -2a + 1$
- Ex:  $a = 1, b = 0$
- Minimize  $||a|| \Rightarrow a = .66, b = -.33$ 
  - Two data on the margin; constraints “tight”




+

# Machine Learning and Data Mining

## Support Vector Machines: Lagrangian and Dual

Prof. Alexander Ihler



# Lagrangian optimization

- Want to optimize constrained system:

$$\theta = (w, b)$$

$$w^* = \arg \min_{w, b} \underbrace{\sum_j w_j^2}_{f(\theta)} \quad s.t. \quad \underbrace{1 - y^{(i)}(w \cdot x^{(i)} + b)}_{g_i(\theta) \leq 0} \leq 0$$

- Introduce Lagrange multipliers  $\alpha$  (one per constraint)


$$\theta^* = \arg \min_{\theta} \max_{\alpha \geq 0} f(\theta) + \sum_i \alpha_i g_i(\theta)$$

- Can optimize  $\theta, \alpha$  jointly, with a simple constraint set
- Then:
  - $g_i(\theta) \leq 0 : \alpha_i = 0$
  - $g_i(\theta) > 0 : \alpha_i \rightarrow +\infty$
- Any optimum of the original problem is a saddle point of the new
- KKT complementary slackness:  $\alpha_i > 0 \Rightarrow g_i(\theta) = 0$

# Optimization

- Use Lagrange multipliers
  - Enforce inequality constraints

$$w^* = \arg \min_w \max_{\alpha \geq 0} \frac{1}{2} \sum_j w_j^2 + \sum_i \alpha_i (1 - y^{(i)} (w \cdot x^{(i)} + b))$$



Alphas > 0 only on the margin:  
"support vectors"

**Stationary conditions wrt  $w$ :**

$$w^* = \sum_i \alpha_i y^{(i)} x^{(i)}$$

and since any support vector has  $y = wx + b$ ,

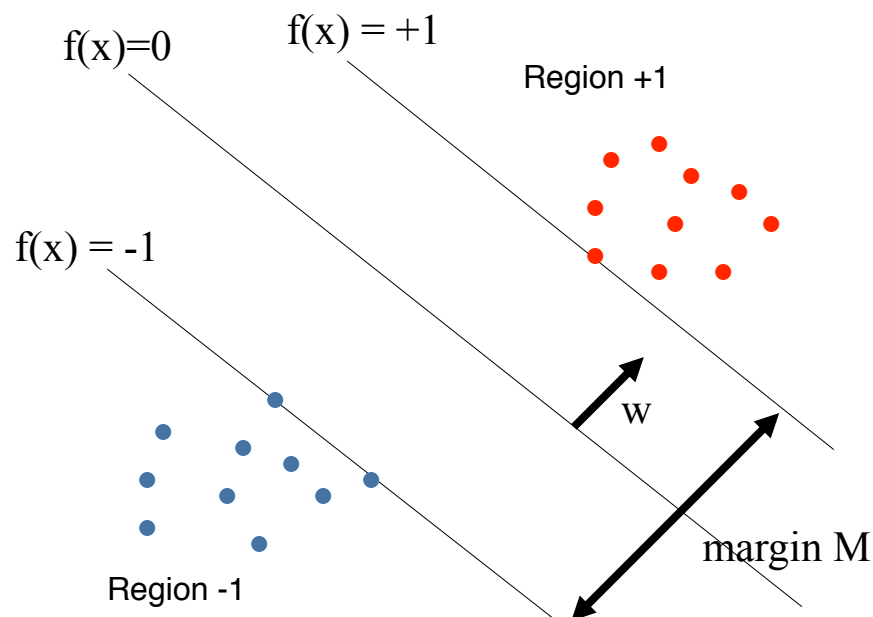
$$b = \frac{1}{N_{sv}} \sum_{i \in SV} (y^{(i)} - w \cdot x^{(i)})$$

# Dual form

- Use Lagrange multipliers
  - Enforce inequality constraints
  - Use solution  $w^*$  to write solely in terms of alphas:

$$\max_{\alpha \geq 0} \sum_i \left[ \alpha_i - \frac{1}{2} \sum_j \alpha_i \alpha_j y^{(i)} y^{(j)} (x^{(i)} \cdot x^{(j)}) \right]$$

$$\text{s.t. } \sum_i \alpha_i y^{(i)} = 0 \quad (\text{since derivative wrt } b = 0)$$



Another quadratic program:  
optimize  $m$  vars with  $1+m$  (simple) constraints  
cost function has  $m^2$  dot products

$$w^* = \sum_i \alpha_i y^{(i)} x^{(i)}$$

$$b = \frac{1}{N_{sv}} \sum_{i \in SV} (y^{(i)} - w \cdot x^{(i)})$$

# Maximum margin classifier

- What if the data are not linearly separable?

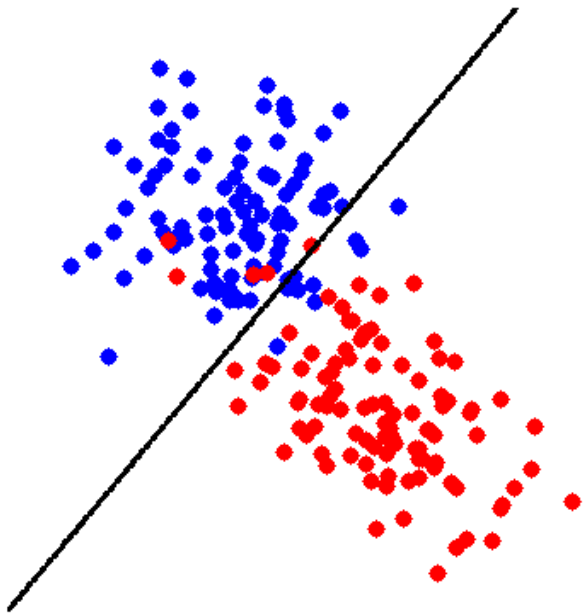
- Want a large “margin”:

$$\min_w \sum_j w_j^2$$

- Want low error:

$$\min_w \sum_i J(y^{(i)}, w \cdot x^{(i)} + b)$$

- “Soft margin” : introduce slack variables for violated constraints



$$w^* = \arg \min_{w, \epsilon} \sum_j w_j^2 + R \sum_i \epsilon^{(i)}$$

*s.t.*

$$y^{(i)} (w^T x^{(i)} + b) \geq +1 - \epsilon^{(i)} \quad (\text{violate margin by } \epsilon)$$

$$\epsilon^{(i)} \geq 0$$

Assigns “cost”  $R$  proportional to distance from margin  
Another quadratic program!



# Maximum margin classifier

- Soft margin optimization:

- For *any* weights  $w$ ,

- we can choose  $\epsilon$  to satisfy constraints

$$w^* = \arg \min_{w, \epsilon} \sum_j w_j^2 + R \sum_i \epsilon^{(i)}$$

$$y^{(i)} (w^T x^{(i)} + b) \geq +1 - \epsilon^{(i)}$$


- Write  $\epsilon^*$  as a function of  $w$  (call this  $J$ ) and optimize directly

- $J$  = distance from the “correct” place

$$J_i = \max[0, 1 - y^{(i)} (w \cdot x^{(i)} + b)]$$

$$w^* = \arg \min_w \frac{1}{R} \sum_j w_j^2 + \sum_i J_i(y^{(i)}, w \cdot x^{(i)} + b)$$

(L2 regularization on the weights)



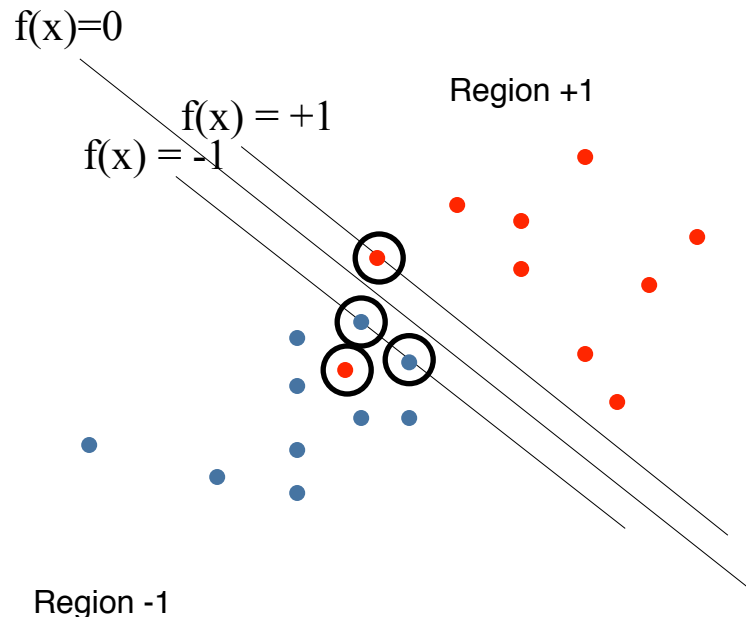
# Dual form

- Soft margin dual:

$$\max_{\underline{0 \leq \alpha \leq R}} \sum_i \alpha_i - \frac{1}{2} \sum_j \alpha_i \alpha_j y^{(i)} y^{(j)} \underbrace{x^{(i)} \cdot x^{(j)}}_{K_{ij}}$$

$K_{ij}$  measures “similarity” of  $x_i$  and  $x_j$  (their dot product)

$$\text{s.t. } \sum_i \alpha_i y^{(i)} = 0$$



Support vectors now data on or past margin...

Prediction:

$$\hat{y} = w^* \cdot x + b = \sum_i \alpha_i y^{(i)} \underbrace{x^{(i)} \cdot x}_{K_{ij}} + b$$

$$w^* = \sum_i \alpha_i y^{(i)} x^{(i)}$$

$b = \dots$  More complicated; can solve e.g. using any  $\alpha \in (0, R)$

# Sequential Minimal Optimization (SMO)

- Out-of-the-box QP solvers not very good for SVMs
- Faster: optimize dual QP coordinate-wise over pairs  $(\alpha_i, \alpha_j)$
- Pick  $\alpha_i, \alpha_j$  s.t.  $\alpha_i$  violates KKT conditions
- Solve constrained QP over just  $(\alpha_i, \alpha_j)$ 
  - Sum constraint  $\Rightarrow$  sum remains constant  $\Rightarrow$  1-D quadratic
  - Upper & lower bounds on alphas

# Multi-class SVMs

- Use standard multi-class linear prediction, 0/1 loss:

$$\hat{y} = f(x; \theta) = \arg \max_y \theta \cdot \Phi(x, y)$$

$$\Phi(x, y) = [ \mathbb{1}[y = 0] \Phi(x) , \mathbb{1}[y = 1] \Phi(x) , \dots ]$$

- Hinge-like loss / slack variable optimization:

$$w^* = \arg \min_{w, b, \epsilon} \sum_j w_j^2 + R \sum_i \epsilon^{(i)}$$

$$w^T \Phi(x^{(i)}, y^{(i)}) - w^T \Phi(x^{(i)}, y) \geq 1 - \epsilon^{(i)} \quad \forall y \neq y^{(i)}$$

- Can introduce class-specific loss function:  $\Delta(y, \hat{y})$

$$w^T \Phi(x^{(i)}, y^{(i)}) - w^T \Phi(x^{(i)}, y) \geq \Delta(y^{(i)}, y) - \epsilon^{(i)} \quad \forall y \neq y^{(i)}$$


- Reduces to earlier form for 0/1 loss:  $\Delta(y, \hat{y}) = \mathbb{1}[y \neq \hat{y}]$
- Again, can optimize as QP (e.g., SMO) or hinge-like loss (e.g., SGD)

+

# Machine Learning and Data Mining


## Support Vector Machines: The Kernel Trick

Prof. Alexander Ihler



# Linear SVMs

- So far, looked at linear SVMs:
  - Expressible as linear weights “w”
  - Linear decision boundary



- Dual optimization for a linear SVM:


$$\max_{0 \leq \alpha \leq R} \sum_i \alpha_i - \frac{1}{2} \sum_j \alpha_i \alpha_j y^{(i)} y^{(j)} (x^{(i)} \cdot x^{(j)}) \quad \text{s.t.} \quad \sum_i \alpha_i y^{(i)} = 0$$

- Depend on pairwise dot products:  $K_{ij} = x^{(i)} \cdot x^{(j)}$ 
  - $K_{ij}$  measures “similarity”, e.g., 0 if orthogonal

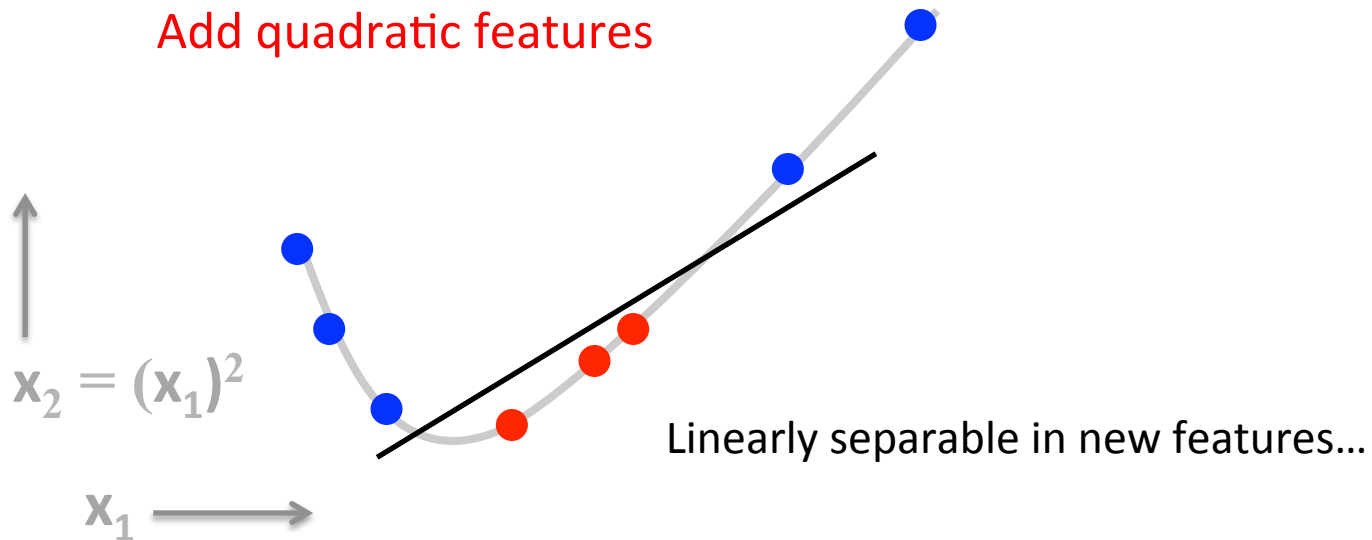
# Adding features

- Linear classifier can't learn some functions

1D example:



Add quadratic features



# Adding features

- Recall: feature function  $\Phi(x)$ 
  - Predict using some transformation of original features

$$\hat{y}(x) = \text{sign}[w \cdot \Phi(x) + b]$$

- Dual form of SVM optimization is:

$$\max_{0 \leq \alpha \leq R} \sum_i \alpha_i - \frac{1}{2} \sum_j \alpha_i \alpha_j y^{(i)} y^{(j)} \Phi(x^{(i)}) \Phi(x^{(j)})^T \quad \text{s.t.} \quad \sum_i \alpha_i y^{(i)} = 0$$

- For example, quadratic (polynomial) features:

$$\Phi(x) = (1 \quad \sqrt{2}x_1 \quad \sqrt{2}x_2 \quad \cdots \quad x_1^2 \quad x_2^2 \quad \cdots \quad \sqrt{2}x_1x_2 \quad \sqrt{2}x_1x_3 \quad \cdots)$$

- Ignore root-2 scaling for now...
- Expands “x” to length  $O(n^2)$



# Implicit features

- Need  $\Phi(x^{(i)})\Phi(x^{(j)})^T$

$$\Phi(x) = (1 \ \sqrt{2}x_1 \ \sqrt{2}x_2 \ \cdots \ x_1^2 \ x_2^2 \ \cdots \ \sqrt{2}x_1x_2 \ \sqrt{2}x_1x_3 \ \cdots)$$

$$\Phi(a) = (1 \ \sqrt{2}a_1 \ \sqrt{2}a_2 \ \cdots \ a_1^2 \ a_2^2 \ \cdots \ \sqrt{2}a_1a_2 \ \sqrt{2}a_1a_3 \ \cdots)$$

$$\Phi(b) = (1 \ \sqrt{2}b_1 \ \sqrt{2}b_2 \ \cdots \ b_1^2 \ b_2^2 \ \cdots \ \sqrt{2}b_1b_2 \ \sqrt{2}b_1b_3 \ \cdots)$$

$$\Phi(a)^T \Phi(b) = 1 + \sum_j 2a_j b_j + \sum_j a_j^2 b_j^2 + \sum_j \sum_{k>j} 2a_j a_k b_j b_k + \dots$$

$$= (1 + \sum_j a_j b_j)^2$$

$$= K(a, b)$$

Can evaluate dot product in  
only  $O(n)$  computations!

# Mercer Kernels

- If  $K(x, x')$  satisfies Mercer's condition:

$$\int_a \int_b K(a, b) g(a) g(b) da db \geq 0$$


For all datasets  $X$ :

$$g^T \cdot K \cdot g \geq 0$$

- Then,  $K(a, b) = \Phi(a) \cdot \Phi(b)$  for some  $\Phi(x)$
- Notably,  $\Phi$  may be hard to calculate
  - May even be infinite dimensional!
  - Only matters that  $K(x, x')$  is easy to compute:
  - Computation always stays  $O(m^2)$

# Common kernel functions

- Some commonly used kernel functions & their shape:
- Polynomial  $K(a, b) = (1 + \sum_j a_j b_j)^d$





# Common kernel functions

- Some commonly used kernel functions & their shape:

- Polynomial  $K(a, b) = (1 + \sum_j a_j b_j)^d$

- Radial Basis Functions

$$K(a, b) = \exp(-(a - b)^2 / 2\sigma^2)$$



# Common kernel functions

- Some commonly used kernel functions & their shape:




- Polynomial  $K(a, b) = (1 + \sum_j a_j b_j)^d$

- Radial Basis Functions

$$K(a, b) = \exp(-(a - b)^2 / 2\sigma^2)$$

- Saturating, sigmoid-like:

$$K(a, b) = \tanh(ca^T b + h)$$



# Common kernel functions

- Some commonly used kernel functions & their shape:

- Polynomial  $K(a, b) = (1 + \sum_j a_j b_j)^d$

- Radial Basis Functions


$$K(a, b) = \exp(-(a - b)^2 / 2\sigma^2)$$

- Saturating, sigmoid-like:

$$K(a, b) = \tanh(ca^T b + h)$$

- Many for special data types:
  - String similarity for text, genetics

- In practice, may not even be Mercer kernels...



# Kernel SVMs

- Linear SVMs
  - Can represent classifier using  $(w,b) = n+1$  parameters
  - Or, represent using support vectors,  $x^{(i)}$
- Kernelized?
  - $K(x,x')$  may correspond to high (infinite?) dimensional  $\Phi(x)$
  - Typically more efficient to remember the SVs
  - “Instance based” – save data, rather than parameters
- Contrast:
  - Linear SVM: identify *features* with linear relationship to target
  - Kernel SVM: identify *similarity measure* between data(Sometimes one may be easier; sometimes the other!)

# Kernel Least-squares Linear Regression

- Recall L2-regularized linear regression:  $\theta = y X (X^T X + \alpha I)^{-1}$

$$\Rightarrow \theta (X^T X + \alpha I) = y X \xrightarrow{\text{Rearranging,}} \alpha \theta = (y - \theta X^T) X$$

Define:

$$r = \frac{1}{\alpha} (y - \theta X^T) \longrightarrow \underline{\theta} = r X$$

$$\alpha r = \underline{y} - \underline{\theta} \underline{X}^T = \underline{y} - r X X^T$$

Gram matrix:  $m \times m$ ,

$$K_{ij} = \langle x^{(i)}, x^{(j)} \rangle$$

Rearrange & solve for  $r$ :

$$r = (X X^T + \alpha I)^{-1} y = (\mathbf{K} + \alpha I)^{-1} y$$

Linear prediction:

$$\tilde{y} = \langle \theta, \tilde{x} \rangle = r X (\tilde{x})^T = \sum_j r_j \langle x^{(j)}, \tilde{x} \rangle = \sum_j r_j \mathbf{K}(x^{(j)}, \tilde{x})$$

Now just replace  $K(x, x')$  with your desired kernel function!



# Summary

---

- Support vector machines
- “Large margin” for separable data
  - Primal QP: maximize margin subject to linear constraints
  - Lagrangian optimization simplifies constraints
  - Dual QP:  $m$  variables; involves  $m^2$  dot product
- “Soft margin” for non-separable data
  - Primal form: regularized hinge loss
  - Dual form:  $m$ -dimensional QP
- Kernels
  - Dual form involves only pairwise similarity
  - Mercer kernels: dot products in implicit high-dimensional space